

PS_4

NOTE: ANSWERS IN RED

Part 1

- 1) Data set created through Weka. Here is the output, although the csv is also included in the submission:

The screenshot shows the Weka Explorer interface with the 'Classify' tab selected. The classifier chosen is 'IBk -K 1 -W 0 -A "weka.core.neighboursearch.LinearNNSearch -A "weka.core.EuclideanDistance -R first-last"'. The test options are set to 'Cross-validation' with 'Folds' set to 10. The classifier output is displayed in the main window, showing stratified cross-validation results and a detailed accuracy by class table.

Classifier output

=== Stratified cross-validation ===
=== Summary ===

Correctly Classified Instances	629	62.9 %
Incorrectly Classified Instances	371	37.1 %
Kappa statistic	0.4501	
Mean absolute error	0.0421	
Root mean squared error	0.181	
Relative absolute error	59.6518 %	
Root relative squared error	96.6971 %	
Total Number of Instances	1000	

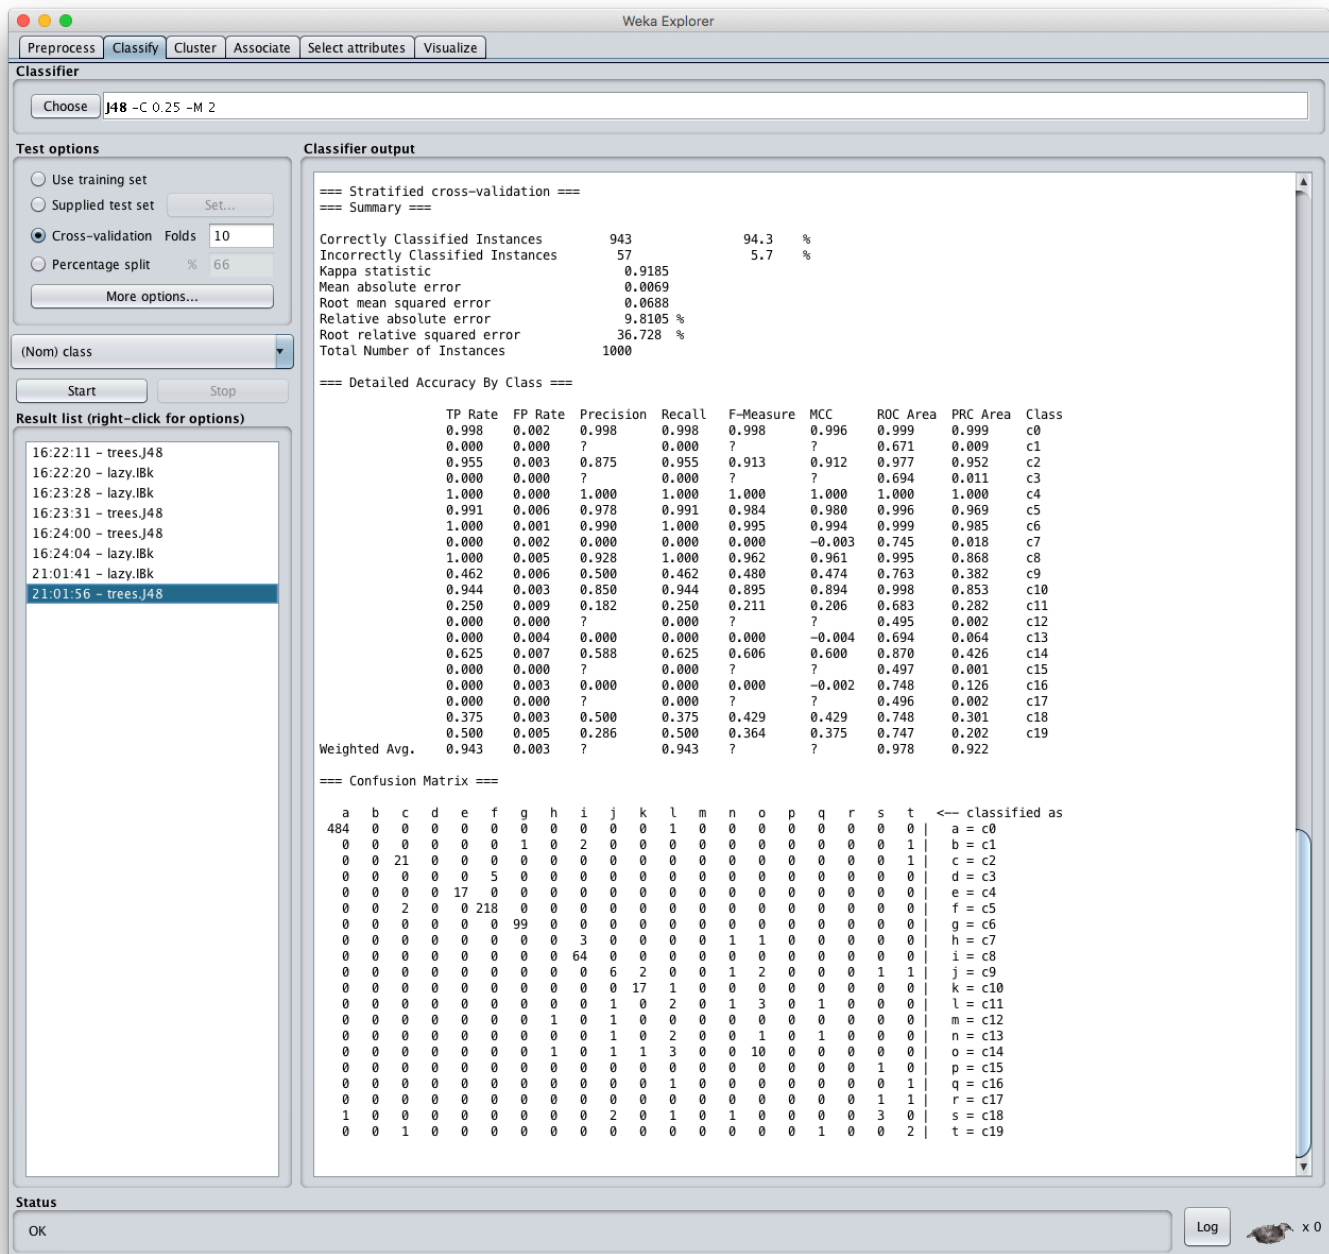
=== Detailed Accuracy By Class ===

	TP Rate	FP Rate	Precision	Recall	F-Measure	MCC	ROC Area	PRC Area	Class
0.876	0.250	0.767	0.876	0.818	0.629	0.874	0.835	c0	
0.000	0.005	0.000	0.000	0.000	-0.004	0.626	0.005	c1	
0.182	0.022	0.154	0.182	0.167	0.147	0.599	0.053	c2	
0.000	0.004	0.000	0.000	0.000	-0.004	0.624	0.038	c3	
0.353	0.016	0.273	0.353	0.308	0.297	0.771	0.193	c4	
0.568	0.090	0.641	0.568	0.602	0.500	0.823	0.538	c5	
0.495	0.047	0.538	0.495	0.516	0.466	0.771	0.332	c6	
0.000	0.002	0.000	0.000	0.000	-0.003	0.487	0.005	c7	
0.219	0.037	0.286	0.219	0.248	0.206	0.726	0.174	c8	
0.077	0.004	0.200	0.077	0.111	0.117	0.548	0.028	c9	
0.111	0.005	0.286	0.111	0.160	0.169	0.578	0.051	c10	
0.250	0.002	0.500	0.250	0.333	0.350	0.796	0.136	c11	
0.000	0.004	0.000	0.000	0.000	-0.003	0.725	0.004	c12	
0.000	0.006	0.000	0.000	0.000	-0.006	0.664	0.028	c13	
0.063	0.013	0.071	0.063	0.067	0.053	0.560	0.025	c14	
0.000	0.000	?	0.000	?	?	0.340	0.001	c15	
0.000	0.002	0.000	0.000	0.000	-0.002	0.860	0.085	c16	
0.000	0.002	0.000	0.000	0.000	-0.002	0.196	0.002	c17	
0.000	0.007	0.000	0.000	0.000	-0.008	0.537	0.020	c18	
0.000	0.001	0.000	0.000	0.000	-0.002	0.426	0.004	c19	
Weighted Avg.	0.629	0.150	?	0.629	?	?	0.808	0.575	

=== Confusion Matrix ===

	a	b	c	d	e	f	g	h	i	j	k	l	m	n	o	p	q	r	s	t	← classified as
425	1	6	0	6	18	12	0	7	2	1	0	0	2	2	0	1	1	1	0		a = c0
1	0	0	0	0	0	0	0	0	2	0	0	0	0	0	0	1	0	0	0	0	b = c1
7	0	4	0	0	6	4	0	0	0	0	0	0	0	0	0	1	0	0	0	0	c = c2
1	0	0	0	0	4	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	d = c3
4	0	0	0	6	5	0	0	0	0	0	0	0	0	1	1	0	0	0	0	0	e = c4
56	1	5	3	5	125	11	0	8	0	1	0	0	1	2	0	0	1	1	0		f = c5
17	0	7	1	1	16	49	0	6	0	1	1	0	0	0	0	0	0	0	0		g = c6
2	0	1	0	0	0	1	0	1	0	0	0	0	0	0	0	0	0	0	0		h = c7
18	2	0	0	2	10	8	2	14	1	0	1	1	1	1	0	0	0	3	0		i = c8
7	0	0	0	0	1	0	0	1	1	0	0	2	0	0	0	0	0	1	0		j = c9
4	0	2	0	0	3	2	0	3	0	2	0	0	0	1	0	0	0	1	0		k = c10
1	0	0	0	0	1	1	0	1	0	0	2	0	0	2	0	0	0	0	0		l = c11
0	0	0	0	0	0	0	0	1	1	0	0	0	0	0	0	0	0	0	0		m = c12
1	0	0	0	1	0	0	0	2	0	0	0	0	0	1	0	0	0	0	0		n = c13
5	1	1	0	1	4	1	0	0	0	0	0	0	1	1	1	0	0	0	0		o = c14
1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0		p = c15
1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1		q = c16
1	0	0	0	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0		r = c17
2	0	0	0	0	1	0	0	3	0	1	0	0	0	1	0	0	0	0	0		s = c18
0	0	0	0	0	1	1	0	0	0	1	0	0	0	0	0	1	0	0	0		t = c19

Nearest Neighbor w/ accuracy 62.9%



Decision Tree w/ accuracy 94.3%

- Absolute difference = $|94.3\% - 62.9\%| = 31.4\%$
- I generated the dataset through Weka itself. I spent a long time generating the data through Excel sheets but realized that wasn't necessary. Through Weka I generated randomly produced data by producing a decision list. In particular, I set the number of attributes to 30, number of classes to 20, kept 1000 examples as requested, and the number of irrelevant attributes to 10. The simple answer for why the classifiers perform so differently is that decision trees are eager learners that first build a classification model on the training dataset and predict a class for a given input vector. Nearest neighbor, however, does not build a classification model and instead learns directly from the training observations. It is coined a lazy learner and must use a distance metric. It lowers in accuracy as we add more dimensions (features).

2) Data set created through Weka. Here is the output, although the csv is also included in the submission:

Weka Explorer

Preprocess Classify Cluster Associate Select attributes Visualize

Classifier: Choose MultilayerPerceptron -L 0.3 -M 0.2 -N 500 -V 0 -S 0 -E 20 -H a

Test options

☐ Use training set

☐ Supplied test set Set...

☒ Cross-validation Folds 10

☐ Percentage split % 66

More options...

(Nom) class

Start Stop

Result list (right-click for options)

- 21:55:26 - bayes.NaiveBayes
- 22:01:01 - bayes.NaiveBayes
- 22:01:07 - functions.MultilayerPerceptron

Classifier output

```

  FALSE      232.0 244.0
  TRUE       269.0 259.0
 [total]     501.0 503.0

Time taken to build model: 0 seconds

=== Stratified cross-validation ===
=== Summary ===

Correctly Classified Instances      544      54.4 %
Incorrectly Classified Instances    456      45.6 %
Kappa statistic                     0.088
Mean absolute error                  0.499
Root mean squared error              0.4998
Relative absolute error              99.7982 %
Root relative squared error          99.9628 %
Total Number of Instances          1000

=== Detailed Accuracy By Class ===
               TP Rate  FP Rate  Precision  Recall   F-Measure  MCC      ROC Area  PRC Area  Cla
               0.537    0.449    0.544     0.537    0.540     0.088    0.537    0.740    TRU
               0.551    0.463    0.544     0.551    0.548     0.088    0.537    0.487    FAL
Weighted Avg.   0.544    0.456    0.544     0.544    0.544     0.088    0.537    0.614

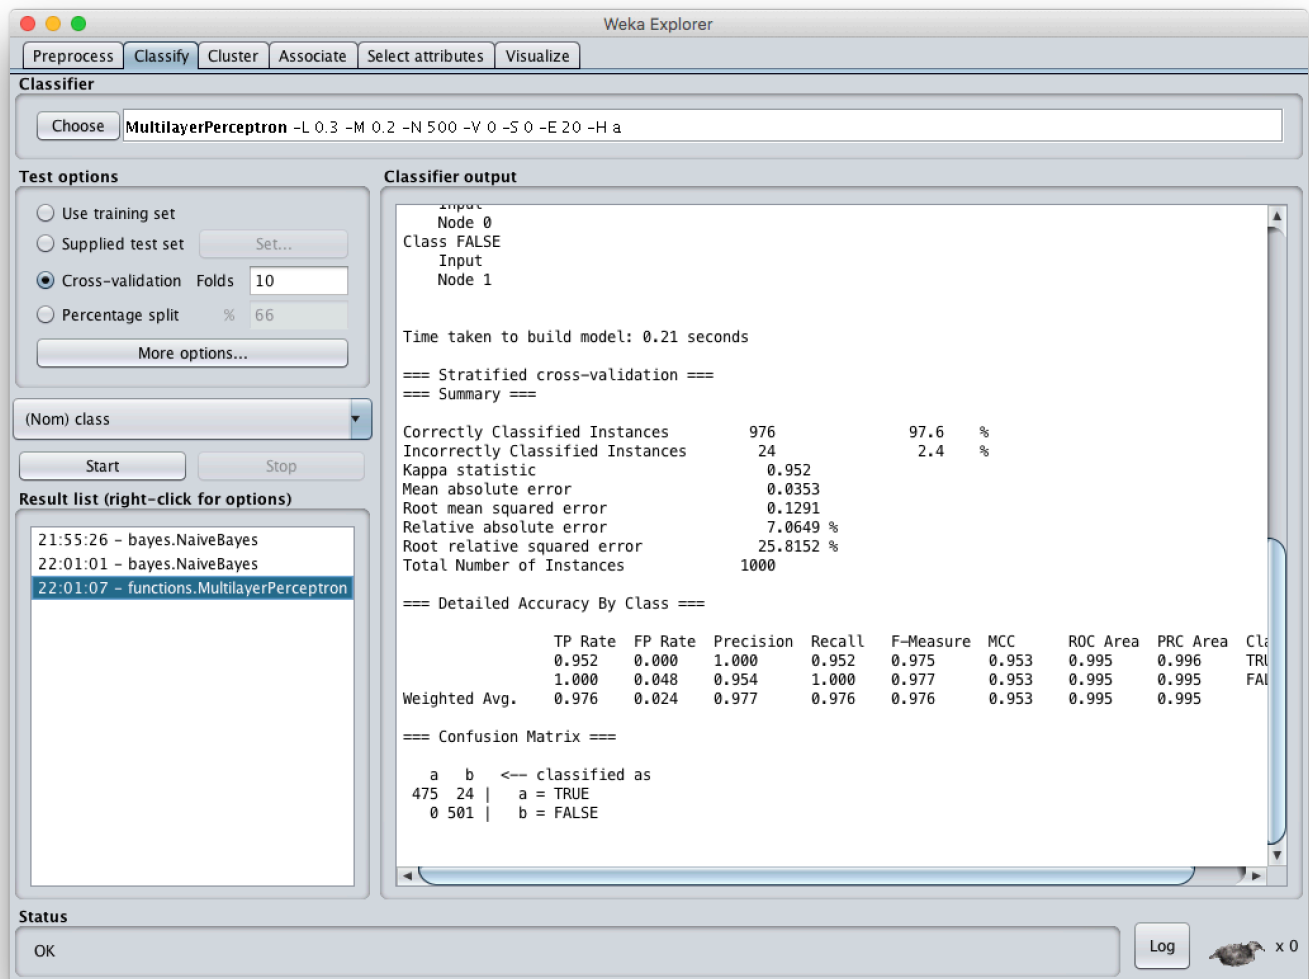
=== Confusion Matrix ===
  a  b  <-- classified as
268 231 | a = TRUE
225 276 | b = FALSE

```

Status

OK Log x 0

Naïve Bayes w/ accuracy 54.4%



MLP w/ accuracy 97.6%

- Absolute difference = $|97.6\% - 54.4\%| = 43.2\%$
- I created the dataset without the help of Weka this time. While my two attributes were randomly true or false, my output was the XOR function. This enabled me to achieve such a low percentage with Naïve Bayes. Naïve Bayes classifiers work by making strong assumptions of conditional between attributes. As mentioned above, I specifically created the data set so that this wasn't the case at all. The MLP model had a better result because it iteratively changes the weights between each successive node to find a stronger correlation and thus classified with better accuracy.

3) Through Jupiter Notebook:

```
In [46]: import numpy as np

from sklearn.naive_bayes import BernoulliNB
from sklearn.linear_model import LogisticRegression

a = BernoulliNB()
b = LogisticRegression(C=5)

In [47]: attributes = np.array([[1,1,1,1],[1,0,0,1],[1,1,0,1]])
output = np.array([0,0,1])

In [48]: a.fit(attributes,output)
b.fit(attributes,output)

Out[48]: LogisticRegression(C=5, class_weight=None, dual=False, fit_intercept=True,
intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
penalty='l2', random_state=None, solver='liblinear', tol=0.0001,
verbose=0, warm_start=False)

In [49]: a.score(attributes,output)

Out[49]: 0.6666666666666666

In [50]: b.score(attributes,output)

Out[50]: 1.0
```

Bernoulli's Model (Naïve Bayes) vs Logistic Regression

- a. For this question, while I worked alone for the most part, I asked Jack Richard for a little guidance (Downey allowed). He suggested I attempt this problem via the Jupiter Notebook. In it I used the Bernoulli Model for Naïve Bayes, and Logistic Regression to achieve a 1.0 and 2/3 training accuracy respectively, as depicted above. In order to achieve a 2/3 training accuracy, it was simpler to get 2 out of 3 example sets to train correctly, while still keeping true to 4 attributes in the matrix labeled `attributes`. Because Naïve Bayes assumes conditional independence, it does not look for any relations between attributes. Logistic Regression can learn a function by a feature. As one of the example sets is classified differently, Naïve Bayes classifies it wrong (2/3). Logistic Regression measures the relationship between the output and one or more independent variables, and probabilistically predicted correctly, thus scoring a 1.0 accuracy.

4) Computing # of parameters

- a. $9 * 3 * 3 * 3 - 1 = 242$ independent parameters
- b. $9 * 3 * 3 * (3 - 1) = 162$ independent parameters
- c. $(9-1)*3 + (3-1)*3 + (3-1)*3 + (3-1) = 38$ independent parameters

5) Attached.

6) Similarity scores:

- a. VGG Representation:
 - i. cat & mj1 = ~15%
 - ii. cat & mj2 = ~14%
 - iii. mj1 & mj2 = ~96%
- b. Pixel Representation:
 - i. cat & mj1 = ~47%
 - ii. cat & mj2 = ~62%

iii. mj1 & mj2 = ~37%

- c. In VGG representation, the most similar pair is MJ1.jpg and MJ2.jpg
In Pixel representation, the most similar pair is Cat.jpg and MJ2.jpg

- 7) The main problem with pixel representation is that it is very sensitive to its position on the image. It does not extract any other features, it simply compares one pixel to another. Pixel representation is considered a “low-level presentation”, and CNNs use an effective supervised learning solution, where the training process uses data that is labeled to close the semantic gap between low-level representations and higher-level ones.
- 8) I actually did something a little different I took the original image, cropped it a little, and returned used that new image to test. Here is the slightly cropped image:



Slightly cropped image of motorcycles. You can compare to the original


- a. VGG Representation provided a better summary of the image than Pixel Representation. They are defined below:
- i. VGG : “A pair of motorcycles parked in a bike slot.”
 - ii. Pixel : “a white red and gray boat some people a bird and some water”
- b. I didn’t try a lot of examples, but for the most part, VGG was better, I think in part because it localizes where an image is in an object. I thought I could break that with cropping, but looks like the prediction worked out just fine for VGG.

9) 2 Images:



Bad VGG Representation : “A tall brown brick building next to a street with traffic”

Pretty sure that we used this example in class. While some elements do exist that are akin to both images (brick building, traffic, etc), the main idea of the images is not there. There is no *tall* brown brick building and there is nothing to represent a street. Still, very close.



Bad Pixel Representation : “Two men that are each trying to catch a Frisbee at the same time”

There are no similar elements to either two men, frisbee, or running around. Basically it’s completely off, but when I looked for an image that did have people playing frisbee, I noticed that the largest similarity was the abundance of green grass in both images. I would imagine that color from both images creates a stronger chance for pixel representation to predict based off of these factors.