

# Linear Regression

EECS 349

slides from Bryan Pardo, Mark Cartwright;  
(also contains ideas and a few images from wikipedia and books by  
Alpaydin, Duda/Hart/ Stork, and Bishop.)

# Outline

---

- ▶ **Announcements**
  - ▶ Homework #2 assigned Wednesday (due Wednesday)
- ▶ **Linear regression**

# Regression Learning

There is a set of possible examples  $X = \{\mathbf{x}_1, \dots, \mathbf{x}_n\}$

Each example is a **vector** of  $k$  **real valued attributes**

$$\mathbf{x}_i = \langle x_{i1}, \dots, x_{ik} \rangle$$

There is a target function that maps  $X$  onto some **real value**  $Y$

$$f : X \rightarrow Y$$

The DATA is a set of tuples  $\langle \text{example}, \text{response value} \rangle$

$$\{\langle \mathbf{x}_1, y_1 \rangle, \dots, \langle \mathbf{x}_n, y_n \rangle\}$$

Find a **hypothesis**  $h$  such that...

$$\forall \mathbf{x}, h(\mathbf{x}) \approx f(\mathbf{x})$$



# Why *Linear* Regression?

---

- ▶ Easily understood/interpretable
- ▶ Well-studied
- ▶ Computationally Efficient



# Linear Regression Assumption

---

- ▶ Response is a linear function of input, plus Gaussian Noise

Observed response  $y = f(\mathbf{x}) + \varepsilon$

Where  $\varepsilon \sim \mathcal{N}(0, \sigma^2)$

# Hypothesis Space

---

- ▶ Each hypothesis characterized by a weight vector  $\mathbf{w}$

$$h(\mathbf{x}) = w_0 + w_1x_1 + w_2x_2 + \dots w_kx_k$$

$$\mathbf{W} = \langle w_0, w_1, \dots, w_k \rangle$$


- ▶ **Goal:** Find a good  $\mathbf{w}$ 
  - ▶ (One that minimizes some error criterion)

# One-dimensional LR

---


- $x$  has 1 attribute  $a$  (predictor variable)
- Hypothesis function is a line:

Example:  $\hat{y} = h(x) = w_0 + w_1 x$



Minimize RSS (sum of squared residuals)

---






# Multivariate Linear Regression

---

$$h(\mathbf{x}) = w_0 + w_1x_1 + w_2x_2 + \dots w_kx_k$$



Create a new 0 dimension with 1 and append it to the beginning of every example vector  $\mathbf{x}_i$

This placeholder corresponds to the offset  $w_0$

$$\mathbf{x}_i = \langle 1, x_{i,1}, x_{i,2}, \dots, x_{i,k} \rangle$$

Format the data as a matrix of examples  $\mathbf{x}$  and a vector of response values  $y$ ...

One training example

$$\mathbf{X} = \begin{bmatrix} 1 & x_{1,1} & \dots & x_{1,k} \\ 1 & x_{2,1} & \dots & x_{2,k} \\ \dots & \dots & \dots & \dots \\ 1 & x_{n,1} & \dots & x_{n,k} \end{bmatrix} \quad \mathbf{y} = \begin{bmatrix} y_1 \\ y_2 \\ \dots \\ y_n \end{bmatrix}$$

# Closed-form solution

---

Our goal is to find the weights of a function....

$$h(\mathbf{x}) = w_0 + w_1 x_1 + w_2 x_2 + \dots w_k x_k$$

...that minimizes the sum of squared residuals:

$$RSS = \sum_i^n (y_i - h(\mathbf{x}_i))^2$$

It turns out that there is a close-form solution to this problem!

$$\mathbf{w} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$



---

$$\begin{aligned}RSS(\mathbf{w}) &= \sum_{i=1}^n (y_i - h(\mathbf{x}_i))^2 \\&= \sum_{i=1}^n (y_i - w_0 - \sum_{j=1}^k x_{ij} w_j)^2 \\&= (\mathbf{y} - \mathbf{X}\mathbf{w})^T (\mathbf{y} - \mathbf{X}\mathbf{w})\end{aligned}$$



$$RSS(\mathbf{w}) = (\mathbf{y} - \mathbf{X}\mathbf{w})^T (\mathbf{y} - \mathbf{X}\mathbf{w})$$

---

$$\frac{\partial RSS}{\partial \mathbf{w}} = -2\mathbf{X}^T (\mathbf{y} - \mathbf{X}\mathbf{w})$$

$$0 = -2\mathbf{X}^T (\mathbf{y} - \mathbf{X}\mathbf{w})$$

$$0 = \mathbf{X}^T (\mathbf{y} - \mathbf{X}\mathbf{w})$$

$$0 = \mathbf{X}^T \mathbf{y} - \mathbf{X}^T \mathbf{X}\mathbf{w}$$

$$\mathbf{X}^T \mathbf{X}\mathbf{w} = \mathbf{X}^T \mathbf{y}$$

$$\mathbf{w} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

---



---

You're familiar with linear regression where the input has  $k$  dimensions.

$$h(\mathbf{x}) = w_0 + w_1x_1 + w_2x_2 + \dots w_kx_k$$

We can use this same machinery to make polynomial regression from a one-dimensional input.....


$$h(x) = w_0 + w_1x + w_2x^2 + \dots w_kx^k$$

---



---

$$h(x) = w_0 + w_1z + w_2z^2 + w_3z^3$$



Parameter estimation (analytically minimizing sum of squared residuals):

One training example

$$\mathbf{X} = \begin{bmatrix} 1 & z_1^1 & \dots & z_1^k \\ 1 & z_2^1 & \dots & z_2^k \\ \dots & \dots & \dots & \dots \\ 1 & z_n^1 & \dots & z_n^k \end{bmatrix} \quad \mathbf{y} = \begin{bmatrix} y_1 \\ y_2 \\ \dots \\ y_n \end{bmatrix}$$

(Note, there is only 1 attribute  $z$  for each training example.  
Those superscripts are powers, since we're doing polynomial regression)

$$\mathbf{w} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$






# Tuning Model Complexity: Example

---


What is your hypothesis for  $f(x)$ ?



# Tuning Model Complexity: Example

---


What is your hypothesis for  $f(x)$ ?



# Tuning Model Complexity: Example


---

What is your hypothesis for  $f(x)$ ?




# Tuning Model Complexity: Example

What is your hypothesis for  $f(x)$ ?




# Tuning Model Complexity: Example

What is your hypothesis for  $f(x)$ ?




# Tuning Model Complexity: Example

What is your hypothesis for  $f(x)$ ?




# Tuning Model Complexity: Example

What happens if we fit to more data?



# Tuning Model Complexity: Example

What happens if we fit to more data?





## Bias and Variance of an Estimator

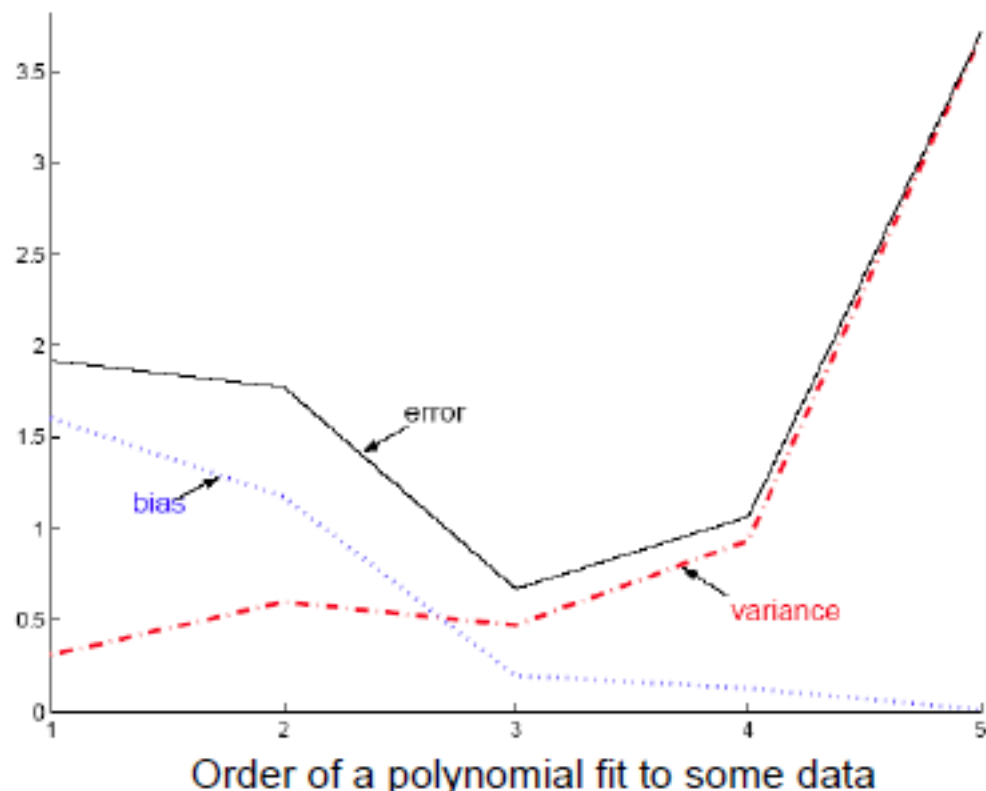
---

- Let  $X$  be a sample from a population specified by a true parameter  $\theta$
- Let  $d=d(X)$  be an estimator for  $\theta$

$$\mathbb{E}[(d - \theta)^2] = \mathbb{E}[(d - \mathbb{E}[d])^2] + (\mathbb{E}[d] - \theta)^2$$

*mean square error*                      *variance*                      *bias<sup>2</sup>*

# Bias and Variance



As we increase complexity, bias decreases (a better fit to data) and variance increases (fit varies more with data)

# Reading

---

- ▶ Chapter 3 from Elements of Statistical Learning
  - ▶ <https://web.stanford.edu/~hastie/ElemStatLearn/>

