

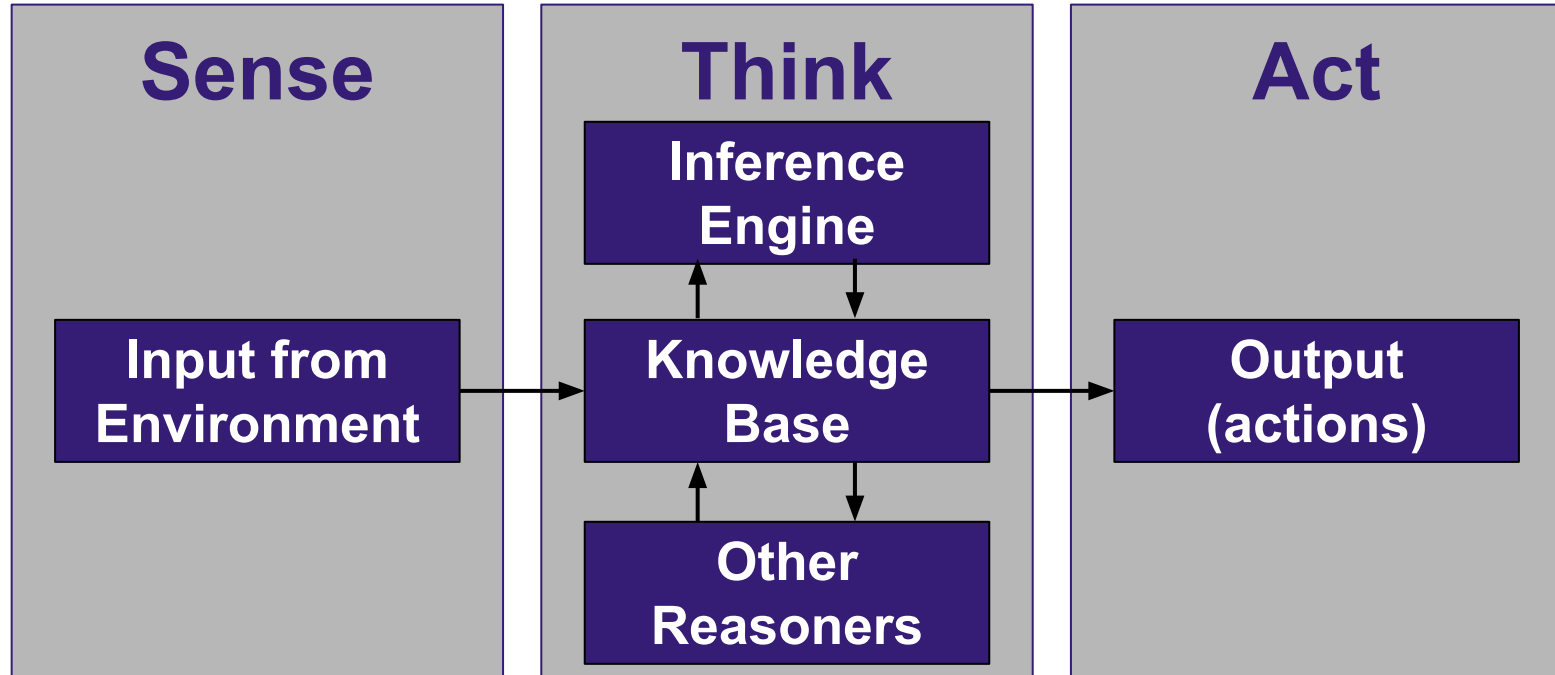
# Informed Search

---


willie

(adapted from slides from Sara Owsley Sood)

# Knowledge-based agent



# Vacuum world knowledge-based agent using search



## 1. **Formulate problem and goal**

2. Search for a sequence of actions that will lead to the goal (the policy)
3. Execute the actions one at a time

Well-defined problem:

(State space)

Initial state

Goal test

Actions/Successor function

Path cost

# Tree Search Algorithm

1. Add the initial state (root) to the <fringe>
2. **Choose a node (curr) to examine from the <fringe>**  
(if there is nothing in <fringe> - FAILURE)
3. Is curr a goal state?  
If so, SOLUTION  
If not, continue
4. Expand curr by applying all possible actions (add the new resulting states to the <fringe>)
5. Go to step 2

# Search algorithm properties

Time (using Big-O) aka: branching factor?

b = ?? prob # of branches

d = shallowest depth

m = max depth of our tree

- approximate the number of nodes generated (not necessarily examined)

Space (using Big-O)

- the max # of nodes stored in memory at any time

Complete

- If a solution exists, will we find it?

Optimal

- If we return a solution, will it be the best/optimal (really just shallowest) solution

# Uninformed search strategies

**Uninformed** search strategies use only the information available in the problem definition

- Breadth-first search
- Depth-first search
- Depth-limited search
- Iterative deepening search


optimal, will do the shortest search

# Breadth-first search

Expand shallowest unexpanded node

Implementation:

- *fringe* is a FIFO queue, i.e., new successors go at end




# Depth-first search

Expand deepest unexpanded node

Implementation:

- *fringe* is a LIFO queue, i.e., put successors at front



space complexity will be ??  $O(b \cdot m)$

time complexity: ??  $O(b^m + 1)$



# Depth-limited search

Depth-first search with depth limit  $L$ ,  
i.e., nodes at depth  $L$  have no successors

search to a certain level.  
space complexity is  $O(b \cdot L)$   
and time complexity is  $O(b^L)$





# Iterative deepening search $L = 0$

Limit = 0




# Iterative deepening search $L = 1$

Limit = 1




# Iterative deepening search $L = 2$

Limit = 2



# Iterative deepening search $L = 3$

Limit = 3



# Properties of iterative deepening search

Space

- $O(bd)$

Complete?

- Yes if there is a solution, it will find it

Optimal?

- Yes only increasing level 1 by 1, we will find the optimal method

# Iterative deepening search

Number of nodes generated in a depth-limited search to depth  $d$  with branching factor  $b$ :

Number of nodes generated in an iterative deepening search to depth  $d$  with branching factor  $b$ :

# Time?

$$L = 0: 1$$

$$L = 1: 1 + b$$

$$L = 2: 1 + b + b^2$$

$$L = 3: 1 + b + b^2 + b^3$$

...

$$L = d: 1 + b + b^2 + b^3 + \dots + b^d$$

Overall:

$$(d+1)b^0 + (d)b^1 + (d-1)b^2 + (d-2)b^3 + \dots + 2b^{d-1} + b^d$$

$$= O(b^d)$$

asymptotically it's more or less the same effort ( $b^d$ )

the cost of the repeat of the lower levels is subsumed by the cost at the highest level



# Properties of iterative deepening search

Time?

$$(d+1)b^0 + d b^1 + (d-1)b^2 + \dots + b^d = O(b^d)$$

Space?

$$O(bd)$$

Optimal?

Yes

Complete?

Yes

# Summary of algorithms

Criterion	Breadth-First	Uniform-Cost	Depth-First	Depth-Limited	Iterative Deepening
Complete?	Yes	Yes	No	No	Yes
Time	$O(b^{d+1})$	$O(b^{\lceil C^*/\epsilon \rceil})$	$O(b^m)$	$O(b^l)$	$O(b^d)$
Space	$O(b^{d+1})$	$O(b^{\lceil C^*/\epsilon \rceil})$	$O(bm)$	$O(bl)$	$O(bd)$
Optimal?	Yes	Yes	No	No	Yes

# Tree Search Algorithm

1. Add the initial state (root) to the <fringe>
2. **Choose a node (curr) to examine from the <fringe>**  
(if there is nothing in <fringe> - FAILURE)
3. Is curr a goal state?  
If so, SOLUTION  
If not, continue
4. Expand curr by applying all possible actions (add the new resulting states to the <fringe>)
5. Go to step 2

# Search trees for 8-puzzle

2	8	3
1	6	4
7		5

Initial state


1	2	3
8		4
7	6	5

Goal

What does the search tree using BFS look like for this problem?

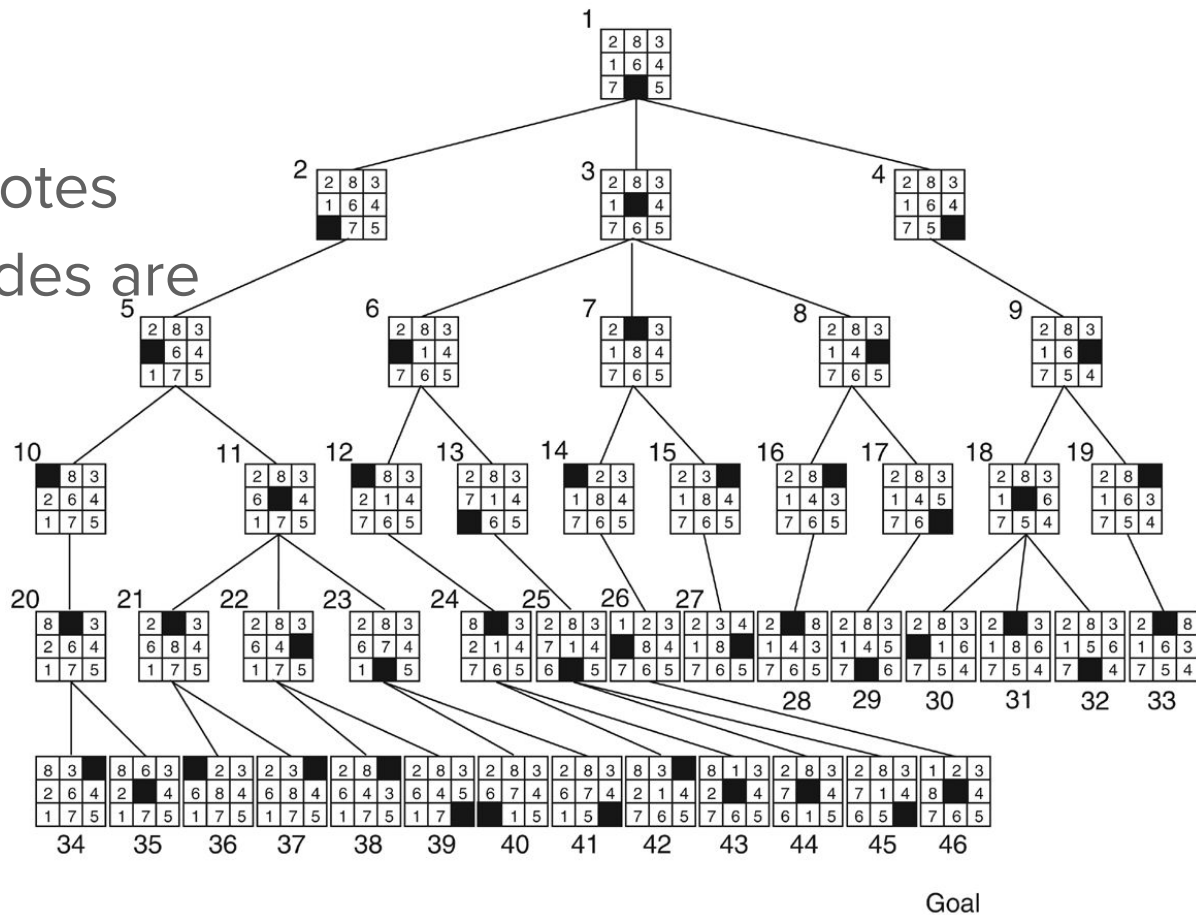
What does the DFS search tree look like?

# Start of BFS search tree



# BFS

Number on node denotes order in which the nodes are visited.




# DFS


Depth bound of 5

Number on node denotes  
the order in which the nodes  
are examined

Pic is missing some nodes  
that were added but not  
examined



Can we do better?



Do we know which of the first choices is best?

Can we make an intelligent choice?



Two heuristics applied to states in the 8-puzzle.

<table border="1"> <tr><td>2</td><td>8</td><td>3</td></tr> <tr><td>1</td><td>6</td><td>4</td></tr> <tr><td></td><td>7</td><td>5</td></tr> </table>	2	8	3	1	6	4		7	5	5	6
2	8	3									
1	6	4									
	7	5									
<table border="1"> <tr><td>2</td><td>8</td><td>3</td></tr> <tr><td>1</td><td></td><td>4</td></tr> <tr><td>7</td><td>6</td><td>5</td></tr> </table>	2	8	3	1		4	7	6	5	3	4
2	8	3									
1		4									
7	6	5									
<table border="1"> <tr><td>2</td><td>8</td><td>3</td></tr> <tr><td>1</td><td>6</td><td>4</td></tr> <tr><td>7</td><td>5</td><td></td></tr> </table>	2	8	3	1	6	4	7	5		5	6
2	8	3									
1	6	4									
7	5										
	Tiles out of place	Sum of distances out of place									

1	2	3
8		4
7	6	5

Goal

# Best-first search

Idea: use an **evaluation function**  $f(n)$  for each node

- estimate of "desirability"
- Expand most desirable unexpanded node

## Implementation:

Order the nodes in fringe in decreasing order of desirability

## Special cases:

- greedy best-first search
- A\* search

Greedy best-first search

Evaluation function

$f(n) = h(n)$  (**h**euristic)

Estimate of cost from  
 $n$  to *goal*

<table><tr><td>2</td><td>8</td><td>3</td></tr><tr><td>1</td><td>6</td><td>4</td></tr><tr><td></td><td>7</td><td>5</td></tr></table>	2	8	3	1	6	4		7	5	<b>5</b>	<b>6</b>
2	8	3									
1	6	4									
	7	5									
<table><tr><td>2</td><td>8</td><td>3</td></tr><tr><td>1</td><td></td><td>4</td></tr><tr><td>7</td><td>6</td><td>5</td></tr></table>	2	8	3	1		4	7	6	5	<b>3</b>	<b>4</b>
2	8	3									
1		4									
7	6	5									
<table><tr><td>2</td><td>8</td><td>3</td></tr><tr><td>1</td><td>6</td><td>4</td></tr><tr><td>7</td><td>5</td><td></td></tr></table>	2	8	3	1	6	4	7	5		<b>5</b>	<b>6</b>
2	8	3									
1	6	4									
7	5										
	Tiles out of place	Sum of distances out of place									

# A\* search

Idea: avoid expanding paths that are already expensive

Evaluation function  $f(n) = g(n) + h(n)$

- $g(n)$  = cost so far to reach  $n$
- $h(n)$  = estimated cost from  $n$  to goal
- $f(n)$  = estimated total cost of path through  $n$  to goal  
(the evaluation of the desirability of  $n$ )

# A\* in 8-puzzle

The evaluation  
function  $f$  applied to  
states in the  
8-puzzle

$g(n) = 0$

Start

2	8	3
1	6	4
7		5

$g(n) = 1$

2	8	3
1	6	4
	7	5

2	8	3
1		4
7	6	5

2	8	3
1	6	4
7	5	

Values of  $f(n)$  for each state,

**6**

**4**

**6**

where:

$$f(n) = g(n) + h(n),$$

$g(n)$  = actual distance from  $n$   
to the start state, and

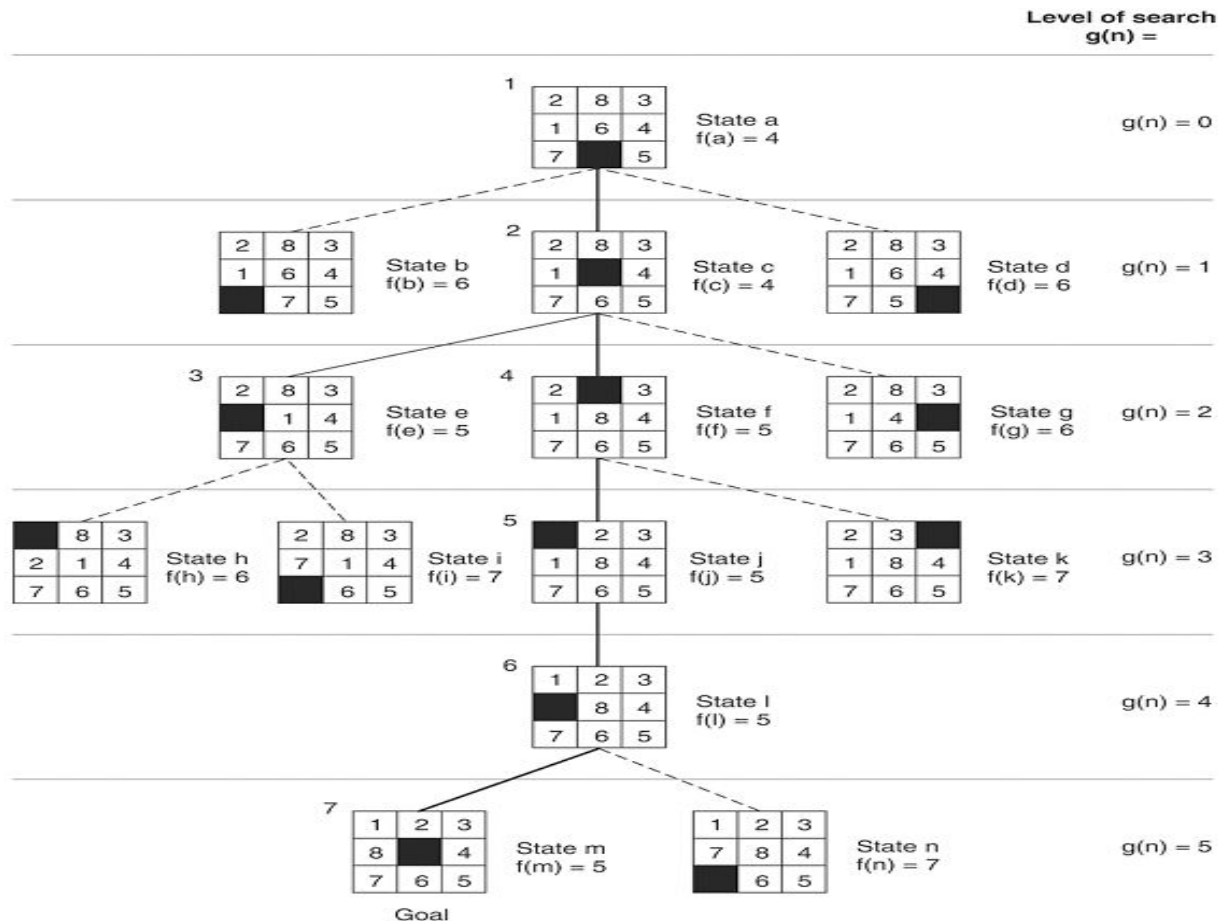
$h(n)$  = number of tiles out of place.

1	2	3
8		4
7	6	5

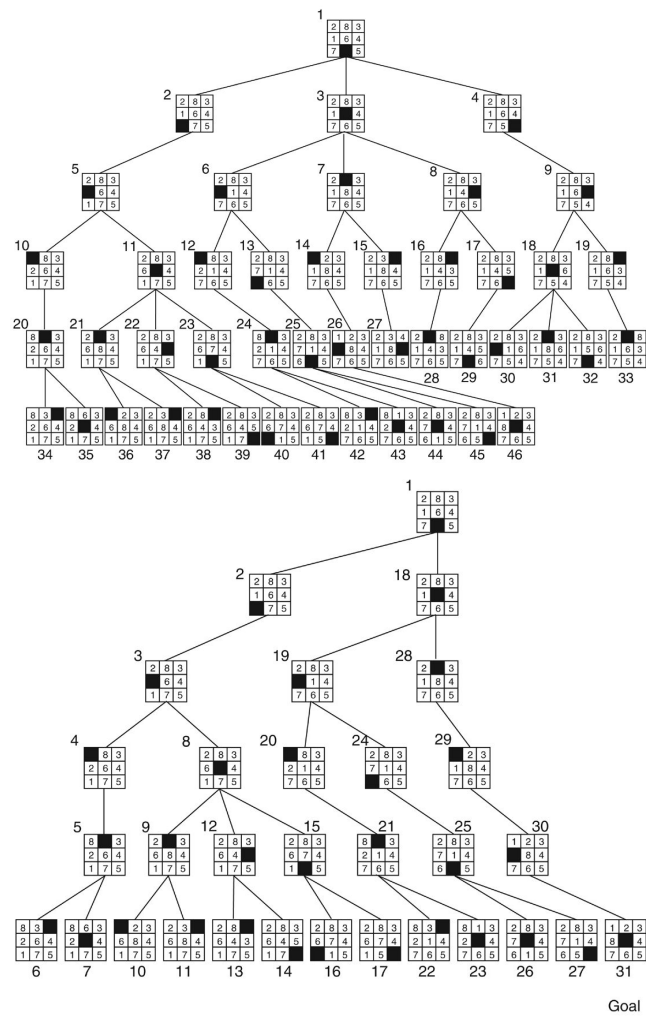
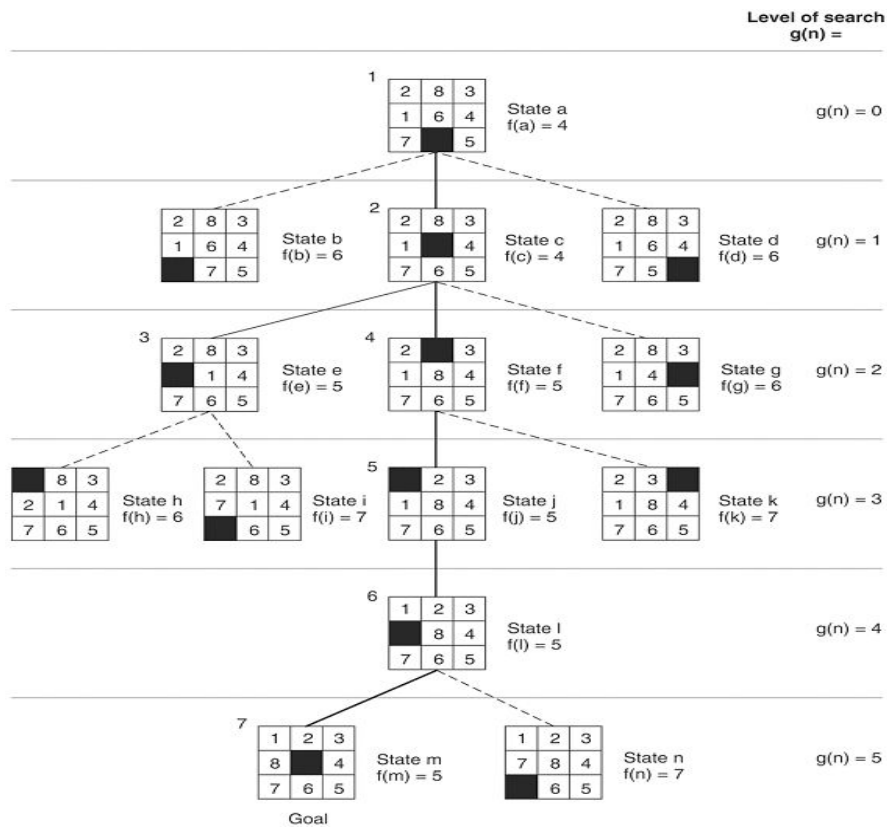
Goal

# State space

Heuristic search for  
8-puzzle is much more  
efficient than BFS and  
DFS




# Compare



# Intelligent order of Expansion?

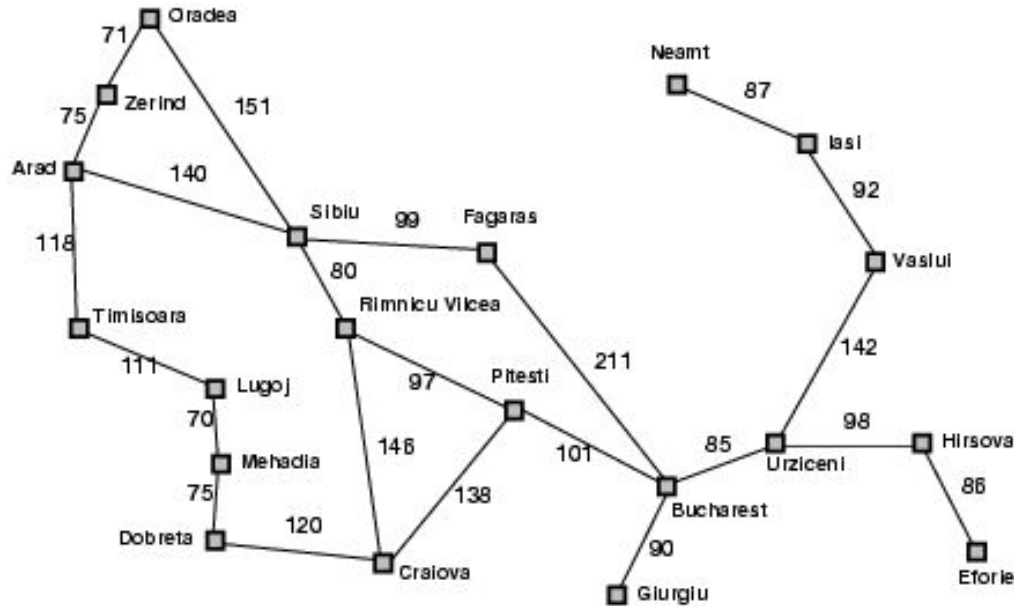
TreeSearch (and GraphSearch) expands its nodes in a given order

Can we be "smart" about the order in which we expand nodes to improve the search?





# Romania with step costs in km



Straight-line distance  
to Bucharest

Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Eforie	161
Fagaras	176
Giurgiu	77
Hirsova	151
Iasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	10
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374

# Greedy best-first search


Evaluation function  $f(n) = h(n)$  (**h**euristic)  
estimate of cost from  $n$  to *goal*

$h_{SLD}(n)$  = straight-line distance from  $n$  to Bucharest


# Greedy best-first search example



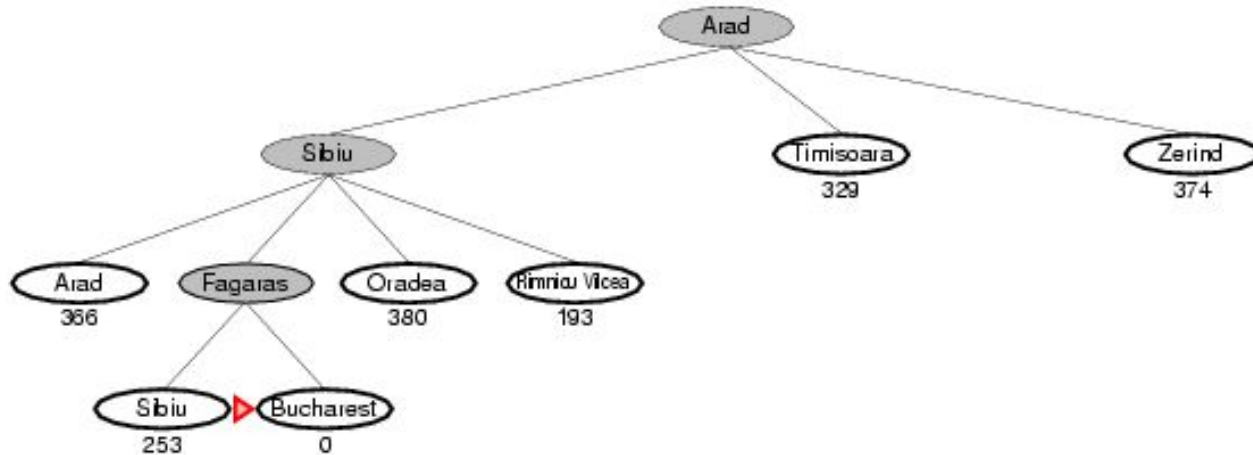
# Greedy best-first search example




# Greedy best-first search example



# Greedy best-first search example



# Problems with greedy best-first search?



Straight-line distance to Bucharest	
Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Eforie	161
Fagaras	176
Giurgiu	77
Hirsova	151
Iasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	10
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374

**SLD to Fagaras**

Neamt – 180

Iasi – 200

Vasliu – 220

Fagaras - 0

# A\* search

Idea: avoid expanding paths that are already expensive

Evaluation function  $f(n) = g(n) + h(n)$

$g(n)$  = cost so far to reach  $n$


$h(n)$  = estimated cost from  $n$  to goal

$f(n)$  = estimated total cost of path through  $n$  to goal


(the evaluation of the desirability of  $n$ )




# A\* search example



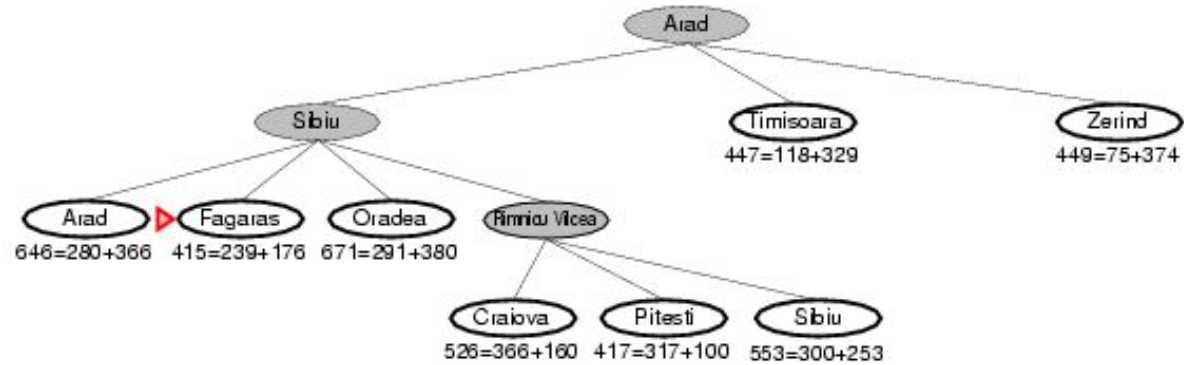
# A\* search example




# A\* search example



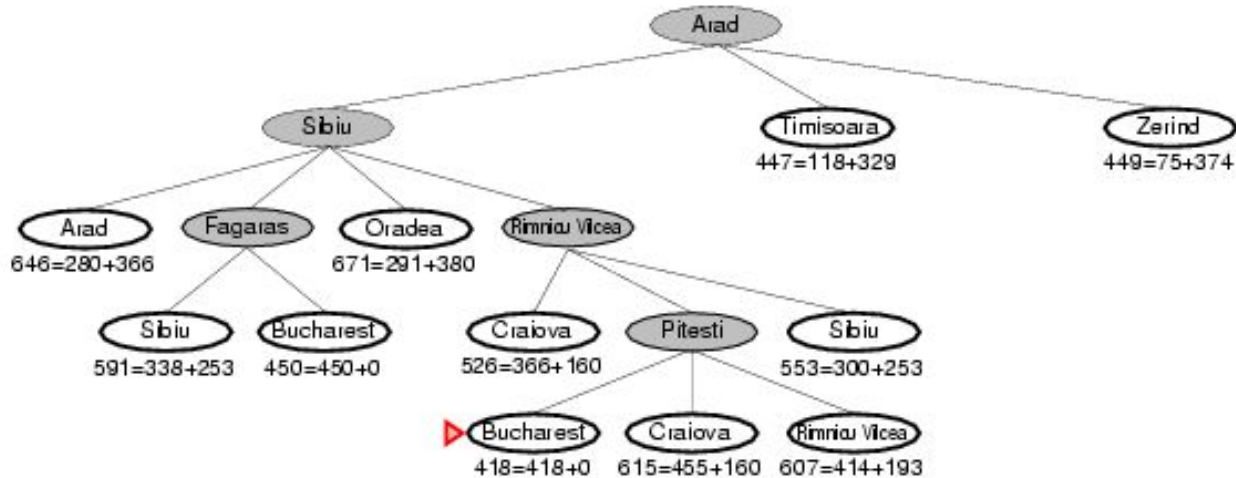
# A\* search example



# A\* search example



# A\* search example



shortest path b/c of the heuristic  
straight line is always nice and admissible heuristic (also consistent)

# Admissible heuristics

A heuristic  $h(n)$  is **admissible** if for every node  $n$ ,

$h(n) \leq h^*(n)$ , where  $h^*(n)$  is the **true** cost to reach the goal state from  $n$ .

An admissible heuristic **never overestimates** the cost to reach the goal, i.e., it is **optimistic**

Example:  $h_{SLD}(n)$  (never overestimates the actual road distance)

# TreeSearch algorithm becomes GraphSearch

1. start with the initial node as *curr*
2. have I been to *curr* before? (is it in CLOSED)
3. is *curr* the goal?
4. if neither, expand *curr* - add children/successors to OPEN, add curr to CLOSED
5. choose a new node *curr* according to the smallest  $f(n)$  & go to step 2