
Adversarial Search


willie

Solving problems with search



Straight-line distant to Bucharest	ce
Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Eforie	161
Fagaras	176
Giurgiu	77
Hirsova	151
Iasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	10
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374

Informed search



For this one example:

A* 7 nodes

BFS 46 nodes

DFS 31 nodes

h2 dominates h1

Averaging over 100 instances of 8-puzzle, at d=12, nodes examined:

```
A*
h1 227
h2 73
IDS 3,644,404
```

Algebra Problems

Solve an algebra problem such as $x^2+3x=0$

- Initial state
- State space
- Successor function
- Goal
- Heuristics

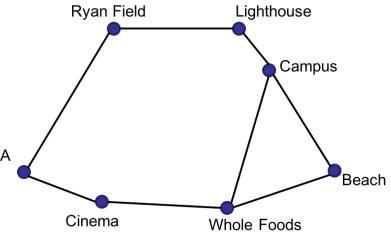
BigByte

Campus —> Cinema? Evaluation function takes in heuristic g(n) = how long it takes you to take the roads

You work at a food delivery start-up called BigByte. The company consists of two people, you and Willie, both alumni of Nerdwestern University located in Heavanston. Of the two, you take care of the software engineering and Willie handles the rest.

BigByte's target client restaurants and ordering customers are all based in Heavanston, which has a handful of landmarks and a heavily fluctuating traffic pattern. To ensure that your delivery fleet works efficiently, Willie asked you to come up with a route planning system.

All major map services are expensive and your start-up begins its services in a week, so you'll have to come up with something yourself and do it fast. Fortunately, Restaurants and Heavanston residents all live close to landmarks, and you only need to tell delivery drivers how to best reach one landmark from another. You have a printed map of all Heavanston roads and data from a traffic survey maintained by University Archive. For you, that means you YWCA only need to write one more piece of software before the launch: an implementation of the A* algorithm that finds the best way to travel between any pair of landmarks in town.



Moving on







"Within ten years a digital computer will be the world's chess champion"

Simon & Newell, 1958

Opponent doesn't want you to win



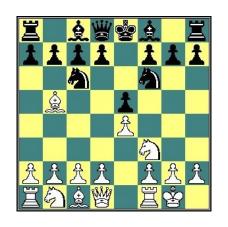


Strategic thinking [?] intelligence

Two-player games have been a key focus of AI as long as computers have been around... Humans and computers have different relative strengths in these games:

humans

good at evaluating the strength of a board for a player



computers

good at looking ahead in the game to find winning combinations of moves

Human Intelligence vs Computer Intelligence:
Evaluation function
General evaluation in humans vs range of options in computers

How humans play games...

An experiment (by deGroot) was performed in which chess positions were shown to

novice and expert players...

- experts could reconstruct these perfectly

- novice players did far worse...



How humans play games...


An experiment (by deGroot) was performed in which chess positions were shown to

novice and expert players...

- experts could reconstruct these perfectly
- novice players did far worse...

Random chess positions (not legal ones) were then shown to the two groups

- experts and novices did just as badly at reconstructing them!



Games' Branching Factors


On average, there are about 35 possible moves that a chess player can make from any board 2

configuration...

Hydra - successor to Deep Blue

Up to **18 ply**





Branching Factor Estimates for different two-player games

Tic-tac-toe 4

Connect Four 4

Checkers 2.8


Othello 10

Chess 35


Go 250

ex. 100 ply = 250^100

How computers play games




2 players:



Define the problem:

- State: state of the board, whose turn
- Successor function: given a player's move, returns a new state
- Goal test: Is game over? Did one player win? Is there a draw?
 - Utility function: Payoff (+1 winning, -1 losing, 0 draw)


How computers play games...




Optimal Strategy

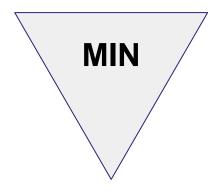
An **Optimal Strategy** is one that is as least as good as any other, no matter what the opponent does

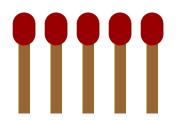
- If there's a way to force the win, it will
- Will only lose if there's no other option



Take 1 or 2 at each turn Goal: take the last match



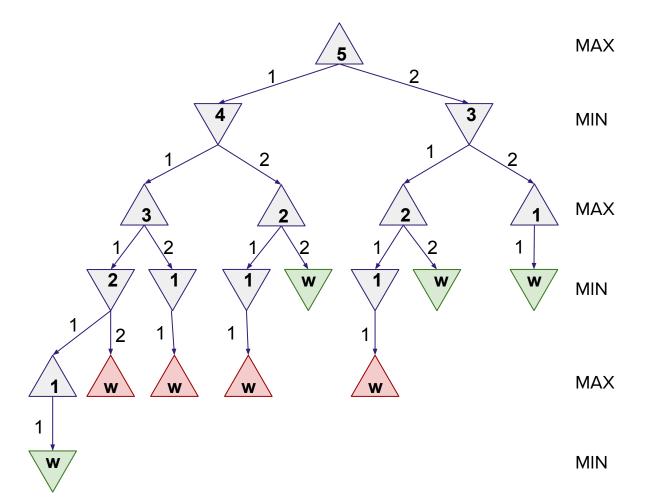


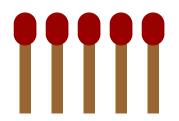


Take 1 or 2 at each turn Goal: take the last match

MAX w = 1


MIN w = -1

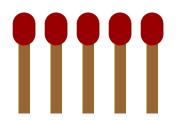




Take 1 or 2 at each turn Goal: take the last match

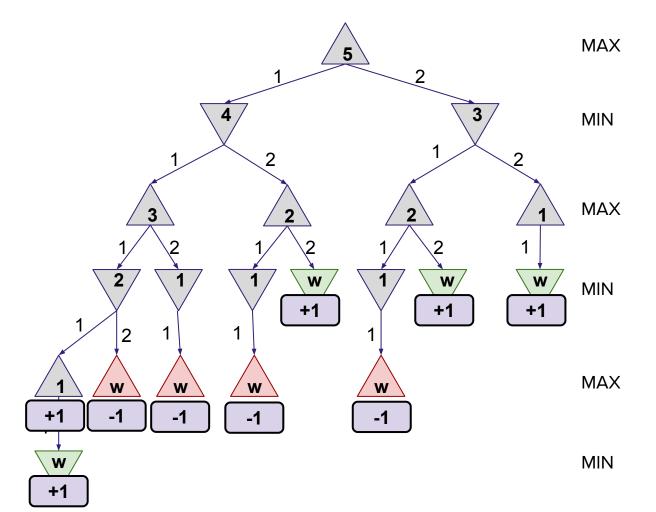
MAX w = 1

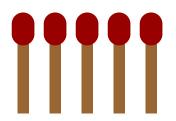




Take 1 or 2 at each turn Goal: take the last match


MAX w = 1

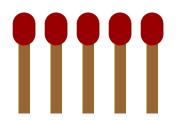




Take 1 or 2 at each turn Goal: take the last match


MAX w = 1






Take 1 or 2 at each turn Goal: take the last match


MAX w = 1

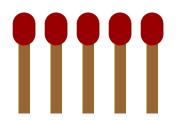




Take 1 or 2 at each turn Goal: take the last match


MAX w = 1






Take 1 or 2 at each turn Goal: take the last match

MAX w = 1



MINIMAX example 2



Minimax: An Optimal Strategy


```
function Minimax-Decision(state) returns an action
   v \leftarrow \text{MAX-VALUE}(state)
   return the action in Successors(state) with value v
function Max-Value(state) returns a utility value
   if Terminal-Test(state) then return Utility(state)
   v \leftarrow -\infty
   for a, s in Successors(state) do
      v \leftarrow \text{Max}(v, \text{Min-Value}(s))
   return v
function Min-Value(state) returns a utility value
   if TERMINAL-TEST(state) then return UTILITY(state)
   v \leftarrow \infty
   for a, s in Successors(state) do
      v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(s))
   return v
```

Minimax Algorithm: An Optimal Strategy

Choose the best move based on the resulting states' MINIMAX-VALUE...

```
MINIMAX-VALUE(n) =
   if n is a terminal state
      then Utility(n)
   else if MAX's turn
      the MAXIMUM MINIMAX-VALUE
      of all possible successors to n
   else if MIN's turn
      the MINIMUM MINIMAX-VALUE
      of all possible successors to n
```

Try it out



It is currently X's turn

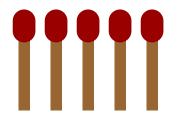
Use MINIMAX to determine what move X should make?

Who will win the game?

https://upload.wikimedia.org/wikipedia/commons/e/e1/Plminmax.gif



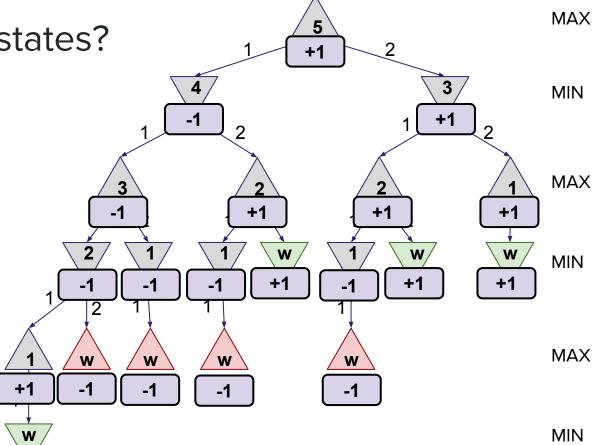
+1



Take 1 or 2 at each turn Goal: take the last match



MIN w = -1



Properties of minimax

For chess, $b \approx 35$, $d \approx 100$ (100 ply) for "reasonable" games

exact solution completely infeasible

Is minimax reasonable for

 Mancala?
 Tic Tac Toe?
 Connect Four?

 B? 6
 B? 4
 B? 4

 D? ?
 D? 9
 D? 20



Resource limits

Suppose we have 100 secs, and can explore 10⁴ nodes/sec

→ can explore 10⁶ nodes per move

Standard approach (Shannon, 1950):

- evaluation function
 estimated desirability of position
- **cutoff test:** e.g., depth limit

Cutting off search

Change:

- if TERMINAL-TEST(state) then return UTILITY(state)
 into
 - if CUTOFF-TEST(state,depth) then return EVAL(state)
- Introduces a fixed-depth limit
 - Is selected so that the amount of time will not exceed what the rules of the game allow
- When cuttoff occurs, the evaluation is performed

Heuristic EVAL

Idea: produce an estimate of the expected utility of the game from a given position.

Performance depends on quality of EVAL.

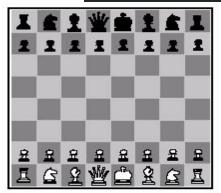
Requirements:

- EVAL should order terminal-nodes in the same way as UTILITY.
- Computation may not take too long.
- For non-terminal states the EVAL should be strongly correlated with the actual chance of winning.

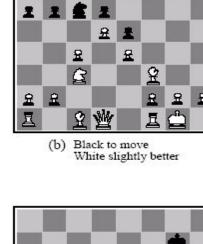
Simple Mancala Heuristic: Goodness of board = # stones in my Mancala minus the number of stones in my opponents.

Heuristic EVAL example

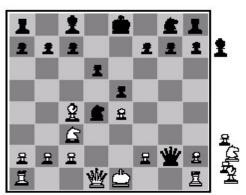
$$Eval(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$



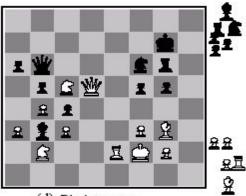
(a) White to move Fairly even



3



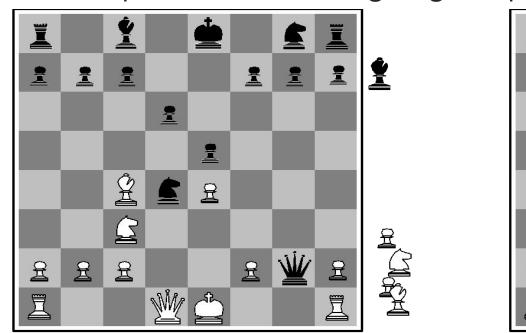
(c) White to move Black winning

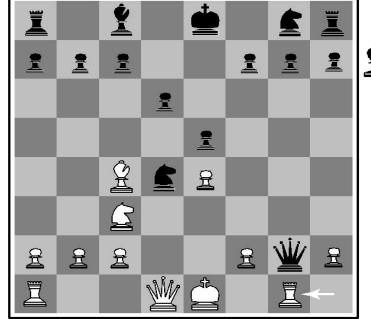


(d) Black to move White about to lose

Heuristic difficulties

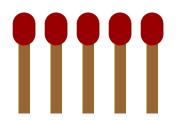
Simple heuristic - weighing the pieces by material value





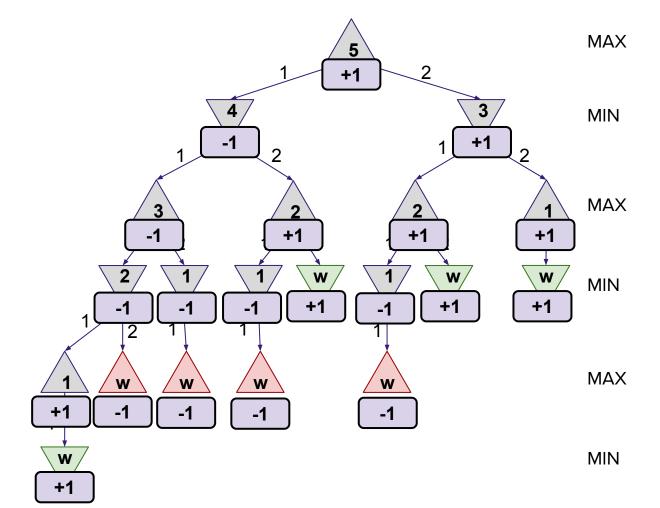
(a) White to move

(b) White to move

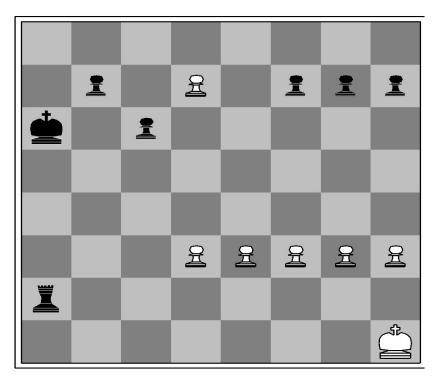


Take 1 or 2 at each turn Goal: take the last match

MAX Wins = 1



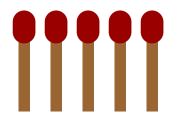
Horizon effect



Fixed depth search thinks it can avoid the queening move

Black to move






Take 1 or 2 at each turn Goal: take the last match



MIN w = -1



Alpha-Beta Pruning

Pruning: eliminate parts of the tree from consideration

Alpha-Beta pruning: prunes away branches that can't possibly influence the final decision

Consider a node n

- If a player has a better choice m (at a parent or further up), then n will never be reached
- So, once we know enough about n by looking at some successors, then we can prune it.